

AI Implications for Business Strategy

3. Machine Learning - Classification



Classification

- The model predicts which of n discrete categories a give input should be classified into
 - We will stick with binary classifications for most of this course
 - For example:
 - Loan risk for a bank customer is high or low
 - Medical test results indicate cancer or no cancer
- When the labeled data is linearly separable (ie the two categories can be separated into two disjoint categories with no errors)
 - Then the problem can be solved with a perceptron
 - We will deal with perceptrons in the Deep Learning section
- When there uncertainty about the classification, it is a *stochastic* classification



Separability







image Credit: http://www.statistics4u.com/fundstat_eng/cc_data_structure.html



The Margin

- We are not interested in the non-linear problem yet
- Margins
 - The margin is the minimum distance from the data points to the separation boundary
 - The margin is negative if there are miss-classified points
 - We use stochastic classification to deal with the cases close to the margin
- We want to assign labels of 1 or 0 to the correctly classified points
 - And a probability of how likely a point is to be correctly classified when it is close to the margin




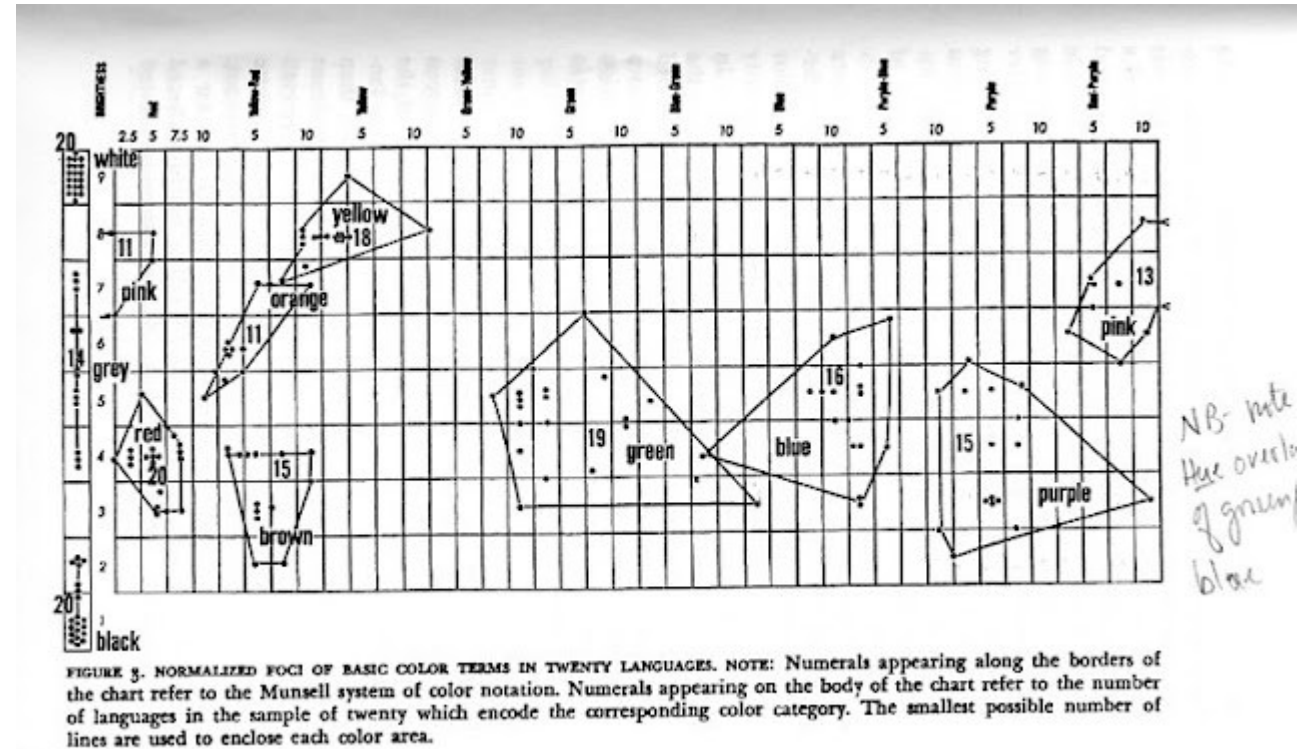
Fuzzy and Crisp Sets

- In crisp sets, membership is binary – classical set theory
 - There is as clear boundary between what is in the set and what is outside
- For fuzzy sets, membership is a gradient
 - Usually determined by how much an data point resembles or is like a prototype
 - This prototype is often called an exemplar or best example of that class
 - An element may have partial membership is several different fuzzy sets
- For Stochastic regression
 - We try to compute a likely hood of a data point belonging to a fuzzy set
 - Done by comparing by developing a probabilistic classification based on features



Color Terms

- There are best examples of the primary colors
- But different languages assign variants to different categories
 - Depends on the amount a similarity to the prototypes
 - The classification rules are language specific
 - The features could be primary colors, hue, saturation, brightness etc.



The Sigmoid Function

- Sigmoid function maps the real number line to the interval $[0,1]$


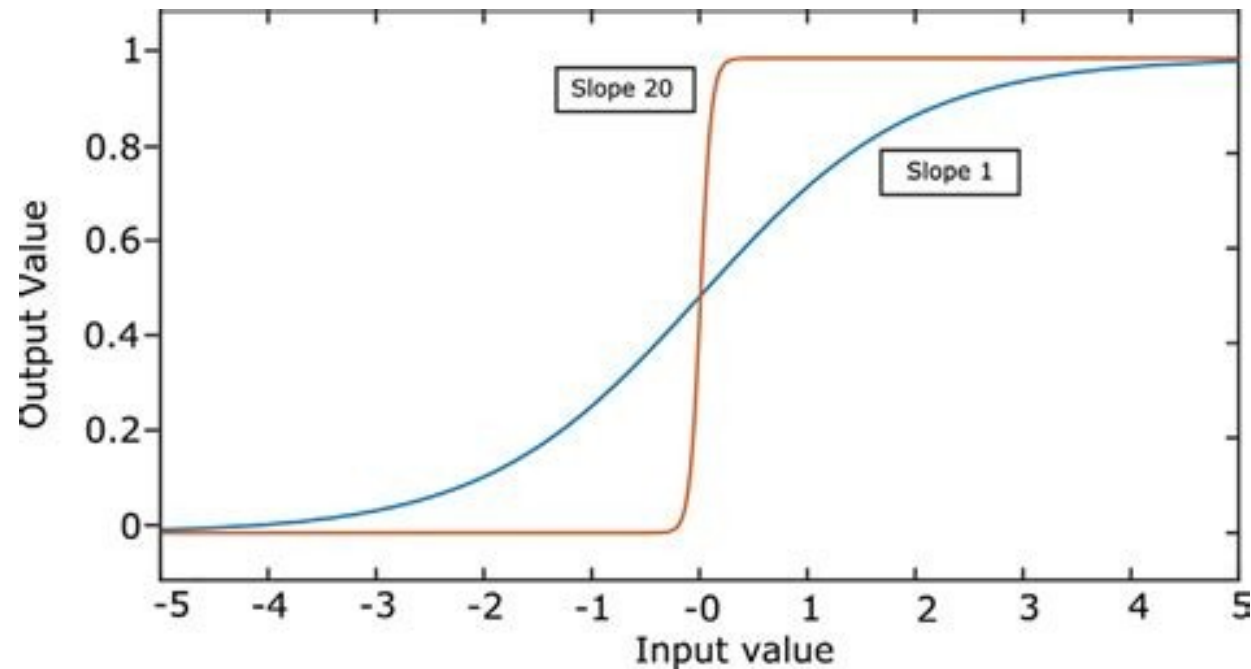


Image Credit: <https://math.stackexchange.com/questions/2437708/sigmoid-circ-sigmoid-sigmoid>



Stochastic Regression

- Finding the right slope of the sigmoid function to describe the probability that data is classified correctly



Sigmoid Alternative

- The hyperbolic tan function is often used to map into the range $[-1,1]$




Image Credit: <https://www.quora.com/What-are-the-benefits-of-a-tanh-activation-function-over-a-standard-sigmoid-activation-function-for-artificial-neural-nets-and-vice-versa>



Other Functions

- Any function could theoretically be used as a classifier
 - Sigmoid and tanh work well with the gradient descent and ML algorithms

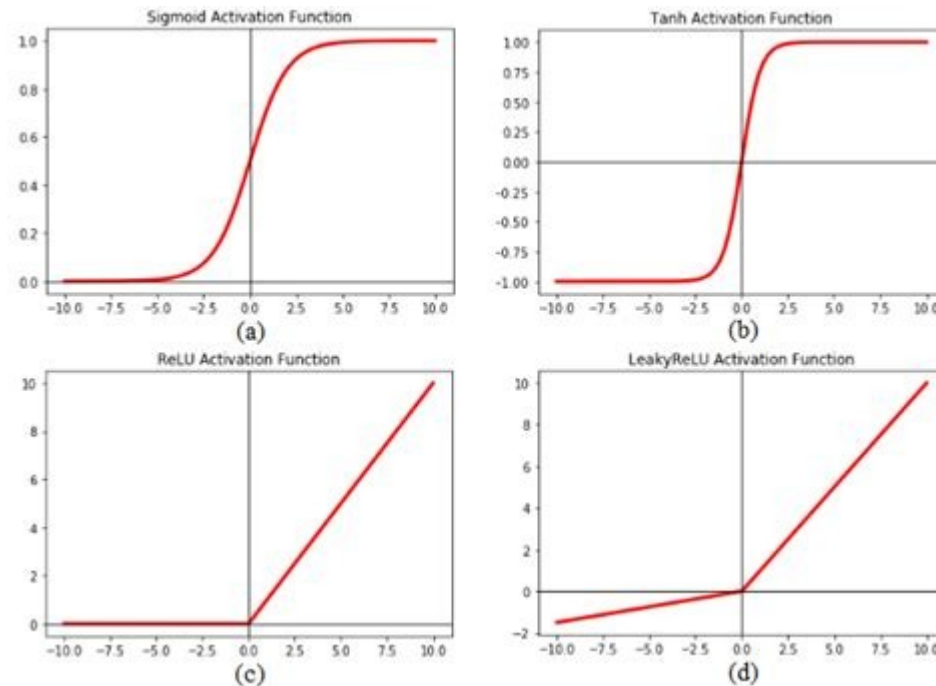


Image Credit: https://www.researchgate.net/figure/Plot-of-different-activation-functions-a-Sigmoid-activation-function-b-Tanh_fig4_339991922



Example: Credit Approval Data

- Given this historical data on credit score and credit card approvals
 - What is the chance some one with score of 700 getting a credit card approved?

Credit Score	Approved?
--------------	-----------

560	No
-----	----

750	Yes
-----	-----

680	Yes
-----	-----

650	No
-----	----

450	No
-----	----

800	Yes
-----	-----

775	Yes
-----	-----

525	No
-----	----

620	No
-----	----


830	Yes
-----	-----

610	Yes
-----	-----




Example: Credit Approval Data

- Graphing, the data is not linearly separable with any vertical line



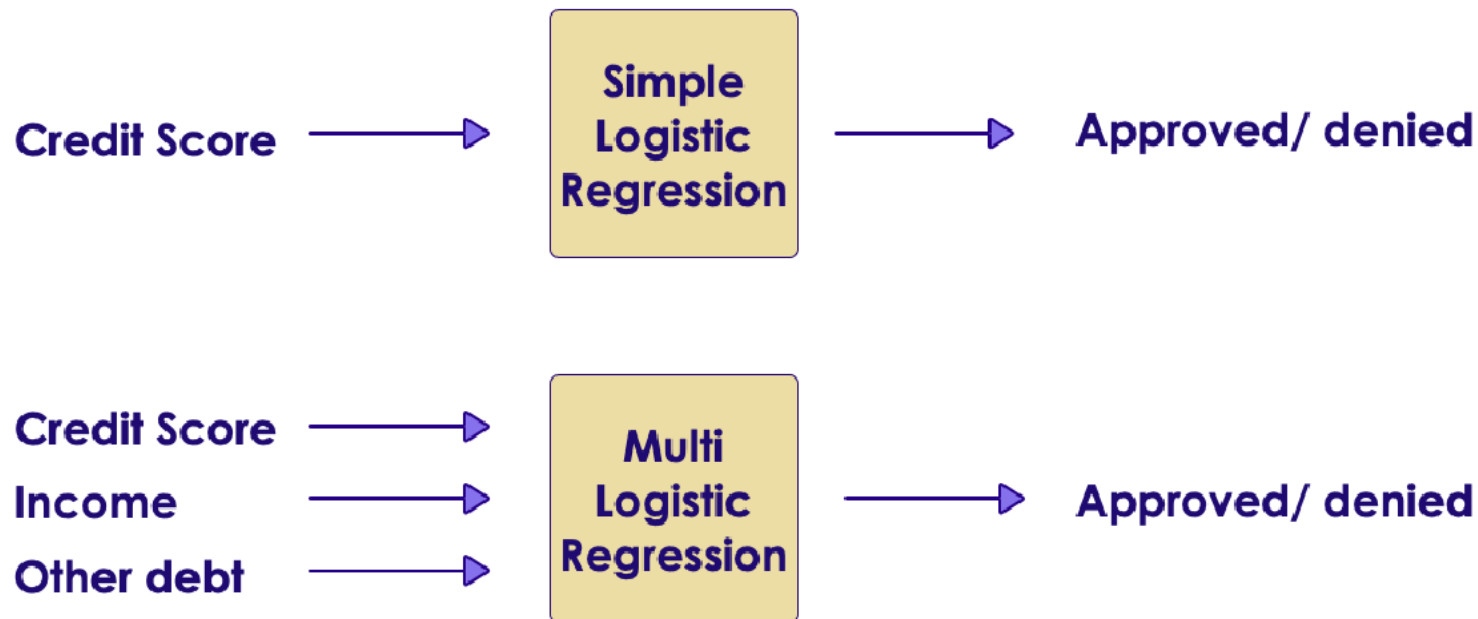
Applying Logistic Regression

- Want to find a stochastic function for data between 600-700
 - Clearly, adjusting the slope of the function will improve performance



Multiple Logistic Regression

- We can have multiple input factors (independent variables) determining a classification as
 - This is called 'multiple logistic regression'



Multiple Logistic Regression




Image Credit: https://www.researchgate.net/figure/Stepwise-multiple-logistic-regression-analyses-of-the-relationships-among-the-mean_fig3_338692557



Preparing Data for Logistic Regression

- Binary Output Variable
 - Logistic Regression predicts probability of an instance belonging to default class
 - These are mapped into 0 or 1 classification based on a threshold value
- Remove noise
 - Remove outliers from input data
- Remove highly correlated inputs to avoid overfitting
- Failure to converge
 - Highly correlated input
 - Data is very sparse (lots of zeros in data)



Multinomial Logistics

- Data output can be a set of discrete categories

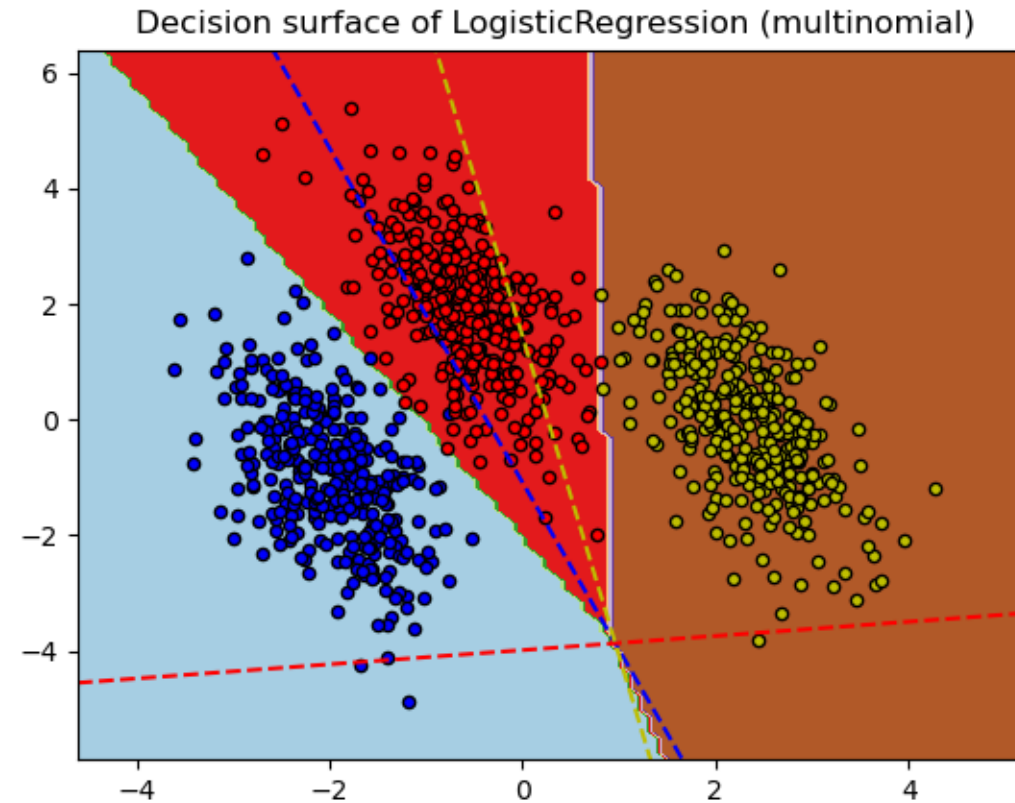


Image Credit: https://w10schools.com/posts/233818_Plot-multinomial-and-One-vs-Rest-Logistic-Regression



Multinomial Logistics

- The final surface is a combination of the individual stochastic functions

Multinomial logistic regression surfaces


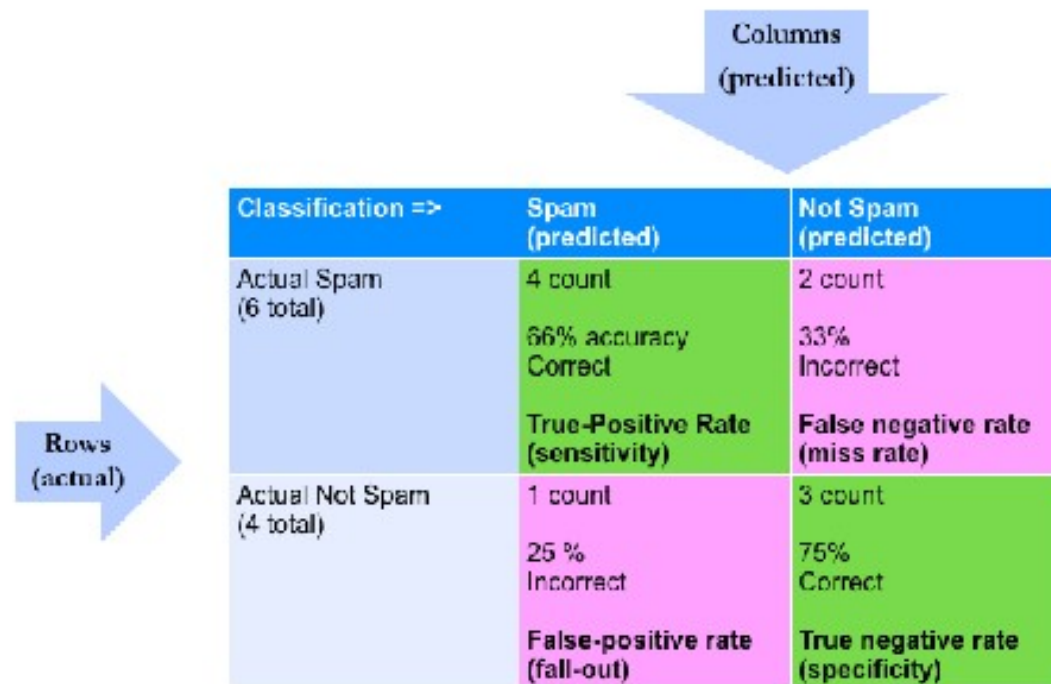


Image Credit: <https://jakejing.github.io/posts/2022/09/blog-post-1/>



Confusion Matrix and ROC Curve

- Confusion Matrix / Error Matrix
 - Binary classifier picks one of two outcomes (spam / not-spam)
 - Say we are classifying 10 emails (6 spam, 4 not-spam)



Classification =>		Columns (predicted)	
		Spam (predicted)	Not Spam (predicted)
Rows (actual)	Actual Spam (6 total)	4 count 66% accuracy Correct True-Positive Rate (sensitivity)	2 count 33% Incorrect False negative rate (miss rate)
	Actual Not Spam (4 total)	1 count 25 % Incorrect False-positive rate (fall-out)	3 count 75% Correct True negative rate (specificity)



Confusion Matrix: More Than 2 Outcomes

		Predicted		
Actual		Cat	Dog	Rabbit
	Cat (8)	5	3	0
	Dog (6)	2	3	1
	Rabbit (13)	0	2	11



Interpreting Confusion Matrix (True/False Positives/Negatives)

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (a cancer diagnostic)	Positive (has cancer)	True positive - Patients who have cancer are correctly identified	False negative <u>Miss rate</u> - A cancer patient is missed - Guilty prisoner was not convicted
	Negative (doesn't have cancer)	False Positive <u>Sensitivity</u> - A healthy patient is flagged incorrectly - False alarm - 'Crying wolf' - Hiring someone who is not qualified	True negative - Patients who do not have cancer are correctly identified



Confusion Matrix: Accuracy / Error Rate

- Accuracy
 - Overall how accurate is the model?
 - $= (TP + TN) / \text{total}$
 - $= (90 + 70) / 200$
 - $= 0.8$ or 80%
- Misclassifications / Error rate
 - How wrong is the model?
 - $= (FP + FN) / \text{total}$
 - $= (10 + 30) / 200$
 - $= 0.2$ or 20%
 - $= 1 - \text{accuracy}$

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90)	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10)	True negative (n = 70)



Confusion Matrix: TPR / FPR

- True Positive Rate (TPR) /Sensitivity / Hit Rate / Recall
 - How often model predicts 'positive' as 'positive' (correctly) ? – actual positive
 - $= TP / (TP + FN)$
 - $= 90 / 120$
 - $= 0.75$ or 75%
- False Positive Rate (FPR)
 - How often model predicts 'negative' as 'positive' (incorrectly) – actual negative = FP / (FP + TN)
 - $= 10 / 80$
 - $= 0.125$ or 12.5%

		Predicted Condition	
		Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5%	True negative (n = 70)



Confusion Matrix: Specificity / Precision / Prevalence

- Specificity

- How often model predicts negative' as negative' (correctly)? – actual no
- $= \text{TN} / (\text{TN} + \text{FP})$
- $= 70 / (70 + 10)$
- $= 0.875$ or 87.5 %
- $= 1 - \text{FPR}$

- Precision / Positive Predictive Value (PPV)

- When model predicts 'positive' how often it is right? – true / predicted positive= $\text{TP} / (\text{TP} + \text{FP})$
- $= 90 / (90 + 10)$
- $= 0.9$ or 90%

	Predicted Condition	
	Predicted Positive	Predicted Negative
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5% True negative (n = 70) Specificity = 87.5%



Confusion Matrix: PPV / Null Error Rate

- Prevalence
 - How often does 'positive' occurs in our sample
 - = actual positive / total
 - = 120 / 200
 - = 0.6 or 60%
- Null Error Rate
 - How often would the model be wrong if it always predicted the majority class?
 - Here our majority = Positive
 - If we always predicted 'positive' we would be wrong 80 times (negative)
 - = 80/200
 - = 40% of time

	Predicted Condition		
	Predicted Positive	Predicted Negative	
Actual condition (n = 200)	Positive (n = 120)	True positive (n = 90) TPR = 75%	False negative (n = 30)
	Negative (n = 80)	False Positive (n = 10) FPR = 12.5%	True negative (n = 70) Specificity = 87.5%



F-Score

- Confusion Matrix: F-Score
 - While precision and recall are very important measures, looking at only one of them will not provide us with the full picture.
 - One way to summarize them is the f-score or f-measure, which is with the harmonic mean of precision and recall
 - $F = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$



Other Measures Associated with Confusion Matrix

Sources: [21][22][23][24][25][26][27][28][29] [view](#) [talk](#) [edit](#)


		Predicted condition			
		Positive (PP)	Negative (PN)	Informedness, bookmaker informedness (BM) $= \text{TPR} + \text{TNR} - 1$	Prevalence threshold (PT) $= \frac{\sqrt{\text{TPR} \times \text{FPR}} - \text{FPR}}{\text{TPR} - \text{FPR}}$
Actual condition	Total population $= P + N$				
	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power $= \frac{\text{TP}}{P} = 1 - \text{FNR}$	False negative rate (FNR), miss rate $= \frac{\text{FN}}{P} = 1 - \text{TPR}$
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection	False positive rate (FPR), probability of false alarm, fall-out $= \frac{\text{FP}}{N} = 1 - \text{TNR}$	True negative rate (TNR), specificity (SPC), selectivity $= \frac{\text{TN}}{N} = 1 - \text{FPR}$
		Prevalence $= \frac{P}{P + N}$	Positive predictive value (PPV), precision $= \frac{\text{TP}}{\text{PP}} = 1 - \text{FDR}$	False omission rate (FOR) $= \frac{\text{FN}}{\text{PN}} = 1 - \text{NPV}$	Positive likelihood ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$
		Accuracy (ACC) $= \frac{\text{TP} + \text{TN}}{P + N}$	False discovery rate (FDR) $= \frac{\text{FP}}{\text{PP}} = 1 - \text{PPV}$	Negative predictive value (NPV) $= \frac{\text{TN}}{\text{PN}}$ $= 1 - \text{FOR}$	Negative likelihood ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$
		Balanced accuracy (BA) $= \frac{\text{TPR} + \text{TNR}}{2}$	F ₁ score $= \frac{2\text{PPV} \times \text{TPR}}{\text{PPV} + \text{TPR}} = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}}$	Fowlkes-Mallows index (FM) $= \sqrt{\text{PPV} \times \text{TPR}}$	Markedness (MK), deltaP (Δp) $= \text{PPV} + \text{NPV} - 1$
				Matthews correlation coefficient (MCC) $= \frac{\sqrt{\text{TPR} \times \text{TNR} \times \text{PPV} \times \text{NPV}} - \sqrt{\text{FNR} \times \text{FPR} \times \text{FOR} \times \text{FDR}}}{1}$	Diagnostic odds ratio (DOR) $= \frac{\text{LR}^+}{\text{LR}^-}$
					Threat score (TS), critical success index (CSI), Jaccard index $= \frac{\text{TP}}{\text{TP} + \text{FN} + \text{FP}}$

Image Credit: https://en.wikipedia.org/wiki/Confusion_matrix




Threshold

- Our spam classifier provides a 'spam probability' for each email
 - Probability is between 0.0 and 1.0 (or 0 to 100%)
 - 1.0 definitely spam
 - 0.0 definitely not spam
 - When an email's 'spam score' is above a certain number we mark it as spam
- This is called 'threshold'



Threshold

- If threshold is lower (say 50%)
 - more emails will be classified as spam (email1, email3)
 - Users will miss emails (as they are in Spam folder)
- If threshold is higher (70%)
 - Fewer emails will be classified as spam (email3)
 - Users will see more spam emails be in Inbox
- We need to find the sweet spot for threshold



Threshold

- In first table our threshold is 0.7
 - 90 emails are correctly predicted as spam
- Next table, our threshold is higher 0.8
 - Only 70 emails are classified as spam Lower TPR

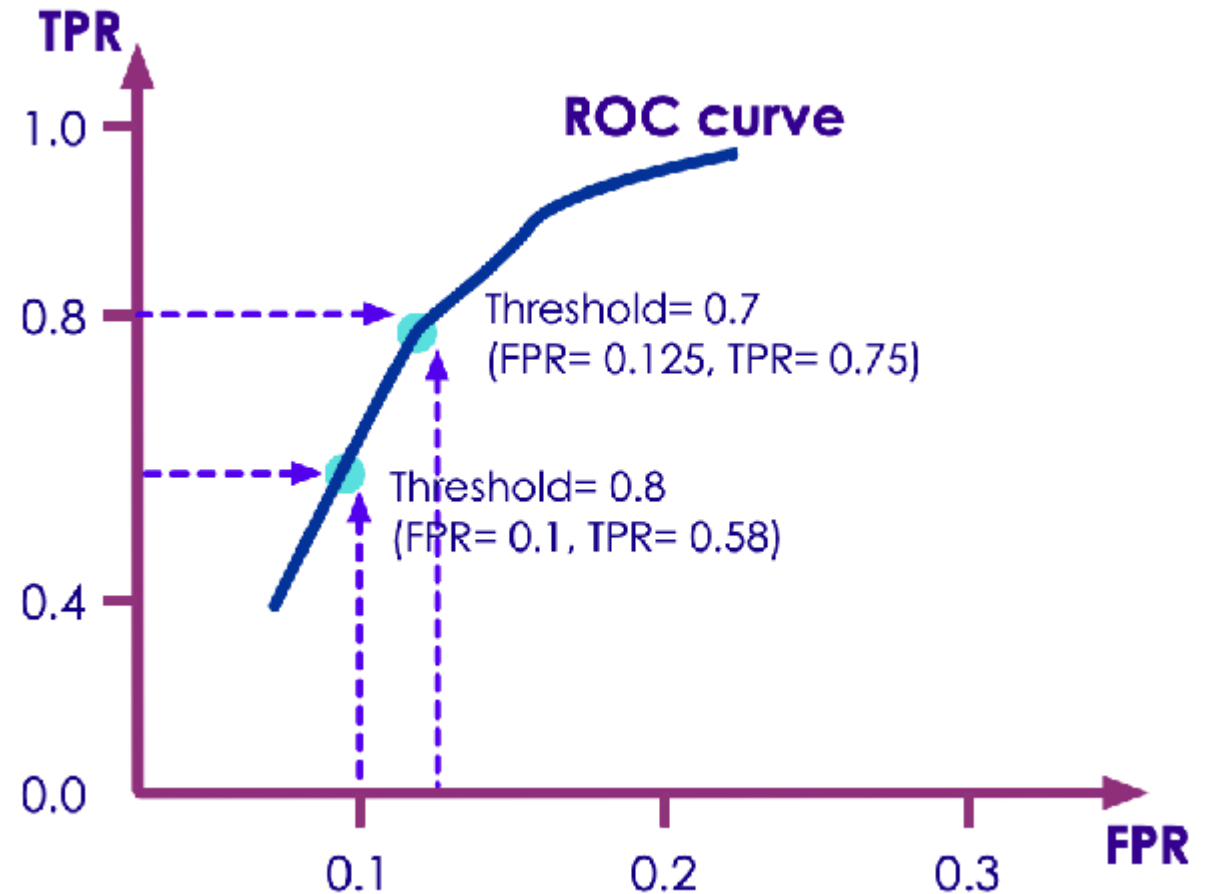
		Predicted Condition	
Threshold = 0.7		Predicted Spam	Predicted Not Spam
Actual condition (total = 200)	Spam (n = 120)	True positive (n = 90) TPR = TP / positive = 90/120 = 75%	False negative (n = 30)
	Not Spam (n = 80)	False Positive (n = 10) FPR = FP / negative = 10/80 = 12.5%	True negative (n = 70)

		Predicted Condition	
Threshold = 0.8 (higher)		Predicted Spam	Predicted Not Spam
Actual condition (total = 200)	Spam (n = 120)	True positive (n = 70) TPR = TP / positive = 70 / 120 = 58.33%	False negative (n = 50)
	Not Spam (n = 80)	False Positive (n = 8) FPR = FP / negative = 8 / 80 = 10%	True negative (n = 72)




ROC Curve

- Receiver Operator Characteristic (ROC) curve
 - a graphical plot used to show the diagnostic ability of binary classifiers.
- Y-axis: True Positive Rate (TPR)
 - Actual=positive, predicted=positive
 - Correct!
- X-axis: False Positive Rate (FPR)
 - Actual=negative, predicted=positive
 - Incorrect!
- $0.0 \leq \text{TPR} \text{ \& \; FPR} \leq 1.0$
- Plot TPR / FPR while varying 'threshold'




ROC Curve

- Shows tradeoff of TPR (sensitivity) vs. FPR (1 - specificity)
- The closer to top-left, the more accurate the model
 - Upper left corner (0,1) = perfect classification!
- The closer to middle line (45 degree) the less accurate the test
- Middle line represents: random classification (50%)



ROC Area Under the Curve

- Measures the percentage of area 'under the curve'
 - AUC is between 0 and 1.0
 - Higher AUC → more accurate the model
- See 3 scenarios below
 - Leftmost is bad (50%), Middle: OK (70%), Rightmost: very good (90%)



Logistic Regression Training


- We use the same process we saw for regression
 - We start with a set of coefficients
 - We have a loss function we can compute
 - We then apply gradient descent to compute a better set of coefficients
- Loss functions for classification
 - Binary Class Entrophy
 - Categorical Crossentropy / Sparse Categorical Crossentropy
 - Negative Log Likelihood
 - Margin Classifier
 - Soft Margin Classifier



Loss Functions Binary Entropy

- Binary Class Entropy
 - Measures the divergence of probability distributions between actual and predicted values

$$E = -\frac{1}{n} \sum_{i=1}^n [y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$



Strengths, Weaknesses, and Parameters

- Strengths
 - Relatively simple to understand
 - Linear algorithm -> Fast learner
 - Works well on high dimensional (100s of features) datasets
 - Very scalable to large data sets
- Weaknesses
 - Can underfit some times
- Parameters
 - Use regularization to minimize overfitting



End of Module

